

Tutorial: Deep Reinforcement Learning

David Silver, Google DeepMind

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models

Reinforcement Learning: $AI = RL$

- ▶ RL is a general-purpose framework for artificial intelligence
- ▶ We seek a single agent which can solve any human-level task
- ▶ The essence of an intelligent agent
- ▶ Powerful RL requires powerful representations

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models

Deep Representations

- ▶ A **deep representation** is a composition of many functions

$$x \xrightarrow{w_1} h_1 \xrightarrow{w_2} \dots \xrightarrow{w_n} h_n \xrightarrow{w_{n+1}} y$$

- ▶ Its gradient can be **backpropagated** by the chain rule

$$\begin{array}{ccccccc} \frac{\partial h_1}{\partial x} & \longleftarrow & \frac{\partial h_2}{\partial h_1} & \longleftarrow & \dots & \longleftarrow & \frac{\partial y}{\partial h_n} \longleftarrow \frac{\partial y}{\partial y} \\ & & \downarrow & & & & \downarrow \\ & & \frac{\partial h_1}{\partial w_1} & & \dots & & \frac{\partial h_n}{\partial w_n} \end{array}$$

$\frac{\partial y}{\partial w_{n+1}}$

Deep Neural Network


A **deep neural network** is typically composed of:

- ▶ Linear transformations


$$h_{k+1} = Wh_k$$

- ▶ Non-linear activation functions


$$h_{k+2} = f(h_{k+1})$$



Step Function



Linear Function




Threshold Logic




Sigmoid Function

Weight Sharing

Recurrent neural network shares weights between time-steps



Convolutional neural network shares weights between local regions



Loss Function

- ▶ A **loss function** $l(y)$ measures goodness of output y , e.g.
 - ▶ Mean-squared error $l(y) = ||y^* - y||^2$
 - ▶ Log likelihood $l(y) = \log \mathbb{P}[y^*|x]$
- ▶ The loss is appended to the forward computation

$$x \xrightarrow{w_1} h_1 \xrightarrow{w_2} \dots \xrightarrow{w_n} h_n \xrightarrow{w_{n+1}} y \longrightarrow l(y)$$

- ▶ Gradient of loss is appended to the backward computation

$$\begin{array}{ccccccc} \frac{\partial h_1}{\partial x} & \longleftarrow & \frac{\partial h_2}{\partial h_1} & \longleftarrow & \dots & \longleftarrow & \frac{\partial y}{\partial h_n} \longleftarrow \frac{\partial l(y)}{\partial y} \\ & & \downarrow & & & & \downarrow \\ & & \frac{\partial h_1}{\partial w_1} & & \dots & & \frac{\partial h_n}{\partial w_n} \end{array}$$

$\frac{\partial y}{\partial w_{n+1}}$

Stochastic Gradient Descent



- ▶ Minimise **expected loss** $\mathcal{L}(w) = \mathbb{E}_x [l(y)]$
- ▶ Follow the **gradient** of $\mathcal{L}(w)$

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E}_x \left[\frac{\partial l(y)}{\partial w} \right] = \mathbb{E}_x \begin{pmatrix} \frac{\partial l(y)}{\partial w^{(1)}} \\ \vdots \\ \frac{\partial l(y)}{\partial w^{(k)}} \end{pmatrix}$$

- ▶ Adjust w in direction of -ve gradient

$$\Delta w = -\frac{\alpha}{2} \alpha \frac{\partial l(y)}{\partial w}$$

where α is a step-size parameter



Deep Supervised Learning

- ▶ Deep neural networks have achieved remarkable success
- ▶ Simple ingredients solve supervised learning problems
 - ▶ Use deep network as a function approximator
 - ▶ Define loss function
 - ▶ Optimise parameters end-to-end by SGD
- ▶ Scales well with memory/data/computation
- ▶ Solves the representation learning problem
- ▶ State-of-the-art for images, audio, language, ...

Deep Supervised Learning

- ▶ Deep neural networks have achieved remarkable success
- ▶ Simple ingredients solve supervised learning problems
 - ▶ Use deep network as a function approximator
 - ▶ Define loss function
 - ▶ Optimise parameters end-to-end by SGD
- ▶ Scales well with memory/data/computation
- ▶ Solves the representation learning problem
- ▶ State-of-the-art for images, audio, language, ...
- ▶ Can we follow the same recipe for RL?

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models

Policies and Value Functions

- **Policy** π is a behaviour function selecting actions given states

$$a = \pi(s)$$

- **Value function** $Q^\pi(s, a)$ is expected total reward from state s and action a under policy π

$$Q^\pi(s, a) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a]$$

“How good is action a in state s ?”

Approaches To Reinforcement Learning

Policy-based RL

- ▶ Search directly for the **optimal policy** π^*
- ▶ This is the policy achieving maximum future reward

Value-based RL

- ▶ Estimate the **optimal value function** $Q^*(s, a)$
- ▶ This is the maximum value achievable under any policy

Approaches To Reinforcement Learning

Policy-based RL

- ▶ Search directly for the **optimal policy** π^*
- ▶ This is the policy achieving maximum future reward

Value-based RL

- ▶ Estimate the **optimal value function** $Q^*(s, a)$
- ▶ This is the maximum value achievable under any policy

Model-based RL

- ▶ Build a transition model of the environment
- ▶ Plan (e.g. by lookahead) using model

Deep Reinforcement Learning

- ▶ Can we apply deep learning to RL?
- ▶ Use deep network to represent value function / policy / model
- ▶ Optimise value function / policy / model **end-to-end**
- ▶ Using stochastic gradient descent

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models

Bellman Equation

- ▶ Bellman expectation equation unrolls value function Q^π

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a] \\ &= \mathbb{E}_{s', a'} [r + \gamma Q^\pi(s', a') \mid s, a] \end{aligned}$$

Bellman Equation

- ▶ **Bellman expectation equation** unrolls value function Q^π

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a] \\ &= \mathbb{E}_{s', a'} [r + \gamma Q^\pi(s', a') \mid s, a] \end{aligned}$$

- ▶ **Bellman optimality equation** unrolls optimal value function Q^*

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Bellman Equation

- ▶ **Bellman expectation equation** unrolls value function Q^π

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a] \\ &= \mathbb{E}_{s', a'} [r + \gamma Q^\pi(s', a') \mid s, a] \end{aligned}$$

- ▶ **Bellman optimality equation** unrolls optimal value function Q^*

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

- ▶ **Policy iteration** algorithms solve Bellman expectation equation

$$Q_{i+1}(s, a) = \mathbb{E}_{s'} [r + \gamma Q_i(s', a') \mid s, a]$$

- ▶ **Value iteration** algorithms solve Bellman optimality equation

$$Q_{i+1}(s, a) = \mathbb{E}_{s', a'} \left[r + \gamma \max_{a'} Q_i(s', a') \mid s, a \right]$$

Policy Iteration with Non-Linear Sarsa

- Represent value function by **Q-network** with weights w

$$Q(s, a, w) \approx Q^\pi(s, a)$$

Policy Iteration with Non-Linear Sarsa

- Represent value function by **Q-network** with weights w

$$Q(s, a, w) \approx Q^\pi(s, a)$$

- Define objective function by mean-squared error in Q-values

$$\mathcal{L}(w) = \mathbb{E} \left[\left(\underbrace{r + \gamma Q(s', a', w)}_{\text{target}} - Q(s, a, w) \right)^2 \right]$$

Policy Iteration with Non-Linear Sarsa

- Represent value function by **Q-network** with weights w

$$Q(s, a, w) \approx Q^\pi(s, a)$$

- Define objective function by mean-squared error in Q-values

$$\mathcal{L}(w) = \mathbb{E} \left[\left(\underbrace{r + \gamma Q(s', a', w)}_{\text{target}} - Q(s, a, w) \right)^2 \right]$$

- Leading to the following **Sarsa** gradient

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E} \left[\left(r + \gamma Q(s', a', w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

- Optimise objective end-to-end by SGD, using $\frac{\partial \mathcal{L}(w)}{\partial w}$

Value Iteration with Non-Linear Q-Learning

- Represent value function by deep Q-network with weights w

$$Q(s, a, w) \approx Q^\pi(s, a)$$

- Define objective function by mean-squared error in Q-values



$$\mathcal{L}(w) = \mathbb{E} \left[\left(\underbrace{r + \gamma \max_{a'} Q(s', a', w)}_{\text{target}} - Q(s, a, w) \right)^2 \right]$$

- Leading to the following **Q-learning** gradient

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

- Optimise objective end-to-end by SGD, using $\frac{\partial \mathcal{L}(w)}{\partial w}$

Example: TD Gammon



Self-Play Non-Linear Sarsa

- ▶ Initialised with random weights
- ▶ Trained by games of self-play
- ▶ Using non-linear Sarsa with afterstate value function

$$Q(s, a, w) = \mathbb{E} [V(s', w)]$$

- ▶ Greedy policy improvement (no exploration)
- ▶ Algorithm converged in practice (not true for other games)


Self-Play Non-Linear Sarsa

- ▶ Initialised with random weights
- ▶ Trained by games of self-play
- ▶ Using non-linear Sarsa with afterstate value function

$$Q(s, a, w) = \mathbb{E} [V(s', w)]$$

- ▶ Greedy policy improvement (no exploration)
- ▶ Algorithm converged in practice (not true for other games)
- ▶ TD Gammon defeated world champion Luigi Villa 7-1 (Tesauro, 1992)

New TD-Gammon Results



Stability Issues with Deep RL

Naive Q-learning **oscillates** or **diverges** with neural nets

1. Data is sequential
 - ▶ Successive samples are correlated, non-iid
2. Policy changes rapidly with slight changes to Q-values
 - ▶ Policy may oscillate
 - ▶ Distribution of data can swing from one extreme to another
3. Scale of rewards and Q-values is unknown
 - ▶ Naive Q-learning gradients can be large
unstable when backpropagated

Deep Q-Networks

DQN provides a stable solution to deep value-based RL

1. Use **experience replay**
 - ▶ Break correlations in data, bring us back to iid setting
 - ▶ Learn from all past policies
 - ▶ Using off-policy Q-learning
2. Freeze **target Q-network**
 - ▶ Avoid oscillations
 - ▶ Break correlations between Q-network and target
3. **Clip** rewards or **normalize** network adaptively to sensible range
 - ▶ Robust gradients

Stable Deep RL (1): Experience Replay

To remove correlations, build data-set from agent's own experience

- ▶ Take action a_t according to ϵ -greedy policy
- ▶ Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory \mathcal{D}
- ▶ Sample random mini-batch of transitions (s, a, r, s') from \mathcal{D}
- ▶ Optimise MSE between Q-network and Q-learning targets, e.g.

$$\mathcal{L}(w) = \mathbb{E}_{s, a, r, s' \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right)^2 \right]$$

Stable Deep RL (2): Fixed Target Q-Network

To avoid oscillations, fix parameters used in Q-learning target

- ▶ Compute Q-learning targets w.r.t. old, fixed parameters w^-

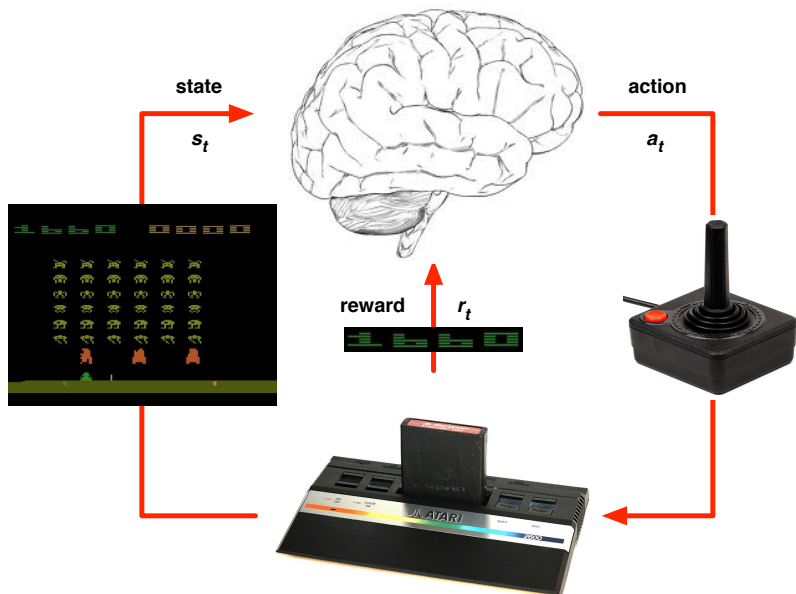
$$r + \gamma \max_{a'} Q(s', a', w^-)$$

- ▶ Optimise MSE between Q-network and Q-learning targets

$$\mathcal{L}(w) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q(s', a', w^-) - Q(s, a, w) \right)^2 \right]$$


- ▶ Periodically update fixed parameters $w^- \leftarrow w$

Reinforcement Learning in Atari




DQN in Atari

- ▶ End-to-end learning of values $Q(s, a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s, a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step



Network architecture and hyperparameters fixed across all games
[Mnih et al.]

DQN Results in Atari



DQN Demo

How much does DQN help?

DQN

	Q-learning	Q-learning + Target Q	Q-learning + Replay	Q-learning + Replay + Target Q
Breakout	3	10	241	317
Enduro	29	142	831	1006
River Raid	1453	2868	4103	7447
Seaquest	276	1003	823	2894
Space Invaders	302	373	826	1089

Stable Deep RL (3): Reward/Value Range

- ▶ DQN clips the rewards to $[-1, +1]$
- ▶ This prevents Q-values from becoming too large
- ▶ Ensures gradients are well-conditioned

Stable Deep RL (3): Reward/Value Range

- ▶ DQN clips the rewards to $[-1, +1]$
- ▶ This prevents Q-values from becoming too large
- ▶ Ensures gradients are well-conditioned
- ▶ Can't tell difference between small and large rewards
- ▶ Better approach: **normalise** network output
- ▶ e.g. via batch normalisation

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models

Policy Gradient for Continuous Actions

- ▶ Represent policy by deep network $a = \pi(s, u)$ with weights u
- ▶ Define objective function as total discounted reward

$$J(u) = \mathbb{E} [r_1 + \gamma r_2 + \gamma^2 r_3 + \dots]$$

- ▶ Optimise objective end-to-end by SGD
- ▶ i.e. Adjust policy parameters u to achieve more reward

Deterministic Policy Gradient

The gradient of the policy is given by

$$\begin{aligned}\frac{\partial J(u)}{\partial u} &= \mathbb{E}_s \left[\frac{\partial Q^\pi(s, a)}{\partial u} \right] \\ &= \mathbb{E}_s \left[\frac{\partial Q^\pi(s, a)}{\partial a} \frac{\partial \pi(s, u)}{\partial u} \right]\end{aligned}$$

Policy gradient is the direction that most improves Q

Deterministic Actor-Critic

Use two networks

- ▶ **Actor** is a policy $\pi(s, u)$ with parameters u

$$s \xrightarrow{u_1} \dots \xrightarrow{u_n} a$$

- ▶ **Critic** is value function $Q(s, a, w)$ with parameters w

$$s, a \xrightarrow{w_1} \dots \xrightarrow{w_n} Q$$

- ▶ Critic provides loss function for actor

$$s \xrightarrow{u_1} \dots \xrightarrow{u_n} a \xrightarrow{w_1} \dots \xrightarrow{w_n} Q$$

- ▶ Gradient backpropagates from critic into actor

$$\frac{\partial a}{\partial u} \longleftarrow \dots \longleftarrow \frac{\partial Q}{\partial a} \longleftarrow \dots \longleftarrow$$

Deterministic Actor-Critic: Learning Rule

- **Critic** estimates value of current policy by Q-learning

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E} \left[\left(r + \gamma Q(s', \pi(s'), w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

- **Actor** updates policy in direction that improves Q

$$\frac{\partial J(u)}{\partial u} = \mathbb{E}_s \left[\frac{\partial Q(s, a, w)}{\partial a} \frac{\partial \pi(s, u)}{\partial u} \right]$$

Deterministic Deep Policy Gradient (DDPG)

- ▶ Naive actor-critic **oscillates** or **diverges** with neural nets
- ▶ DDPG provides a stable solution


Deterministic Deep Policy Gradient (DDPG)

- ▶ Naive actor-critic **oscillates** or **diverges** with neural nets
 - ▶ DDPG provides a stable solution
1. Use **experience replay** for both actor and critic
 2. Freeze **target network** to avoid oscillations

$$\frac{\partial \mathcal{L}(w)}{\partial w} = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[\left(r + \gamma Q(s', \pi(s', u^-), w^-) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$
$$\frac{\partial J(u)}{\partial u} = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[\frac{\partial Q(s, a, w)}{\partial a} \frac{\partial \pi(s, u)}{\partial u} \right]$$

DDPG for Continuous Control

- ▶ End-to-end learning of control policy from raw pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Two separate convnets are used for Q and π
- ▶ Physics are simulated in MuJoCo



DDPG Demo

Outline

Deep Learning

Reinforcement Learning

Deep Value Functions

Deep Policies

Deep Models


Model-Based RL

Learn a **transition model** of the environment

$$p(r, s' \mid s, a)$$

Plan using the transition model

- ▶ e.g. Lookahead using transition model to find optimal actions



Learning Models of the Environment

- ▶ Challenging to plan due to compounding errors
 - ▶ Errors in the transition model compound over the trajectory
 - ▶ Planning trajectories differ from executed trajectories
 - ▶ At end of long, unusual trajectory, rewards are totally wrong

Deep Reinforcement Learning in Go

What if we have a perfect model? e.g. game rules are known

AlphaGo paper:

www.nature.com/articles/nature16961

AlphaGo resources:

deepmind.com/alphago/



Conclusion

- ▶ General, stable and scalable RL is now possible
- ▶ Using deep networks to represent value, policy, model
- ▶ Successful in several tasks
- ▶ Using a variety of deep RL paradigms